TWOSTEP: Multi-agent Task Planning using Classical Planners and Large Language Models

Ishika Singh¹, David Traum¹, Jesse Thomason¹

Abstract—Classical planning formulations like the Planning Domain Definition Language (PDDL) admit action sequences guaranteed to achieve a goal state given an initial state if any are possible. However, reasoning problems defined in PDDL do not capture temporal aspects of action taking, for example that two agents in the domain can execute an action simultaneously if postconditions of each do not interfere with preconditions of the other. A human expert can decompose a goal into largely independent constituent parts and assign each agent to one of these subgoals to take advantage of simultaneous actions for faster execution of plan steps, each using only single agent planning. By contrast, large language models (LLMs) used for directly inferring plan steps do not guarantee execution success, but do leverage commonsense reasoning to assemble action sequences. We combine the strengths of classical planning and LLMs by approximating human intuitions for two-agent planning goal decomposition. We demonstrate that LLM-based goal decomposition leads to faster planning times than solving multiagent PDDL problems directly while simultaneously achieving fewer plan execution steps than a single agent plan alone and preserving execution success. Additionally, we find that LLMbased approximations of subgoals can achieve similar multiagent execution steps than those specified by human experts. Website and resources at glamor-usc.github.io/twostep

I. INTRODUCTION

Symbolic planning problems specified in PDDL tend to explore single agent plans, and changing these domain specifications to enable planning with multiple agents requires a human expert. Unfortunately, multi-agent planning introduces an exponential growth in the search space over possible plans, since each agent can take an action at each timestep. In this paper, we explore whether the commonsense reasoning abilities in LLMs can take advantage of the lexical semantics encoded in the variable, action, and other ontological names of an expert-written planning domain to predict subgoals for individual agents that, when executed together, will achieve a given global goal. We consider both purely symbolic execution domains where transition functions between world states are deterministic and hairier simulation execution domains that introduce stochasticity in environment states and transitions.

Classical planning. Classical or automated task planning algorithms have been widely applied in autonomous space-crafts, military logistics, manufacturing, games, and robotics. The automated STRIPS planner, for example, operated the Shakey robot[1] in 1970. Classical planning algorithms classically work with finite, deterministic, and full state information, admitting guaranteed plan generation when a path from

¹Computer Science Department, University of Southern California {ishikasi, traum, jessetho}@usc.edu

the initial to goal states is possible. Some other frameworks were also shown to be useful for robot planning [2], [3]. Planning domain description language (PDDL) and answer set programming (ASP) are popular specification formats for planning domains [4], [5], [6], [7].

Planning with LLMs. Several works have shown that LLMs can guide agent actions by acting as planners [8], [9], but such stochastic, generative approaches lose the success guarantees of classical planners. Another line of work shows that LLMs plan better with API-based programmatic plan generation [10], [11], [12], which introduces some symbolic structure and constraints but still does not ensure success. Program synthesis for planning has been previously proposed in LEAPS [13], which generates programs using a learned latent space. In recent developments [14], [15], PDDL has been used as a translation medium, with LLMs used to generate either a PDDL plan or goal. Subsequently, a classical planner is used to plan for the PDDL goal. Generating a PDDL goal eliminates the need to track execution state, required when generating a plan using LLMs instead. However, needing to use a classical planner necessitates specification of full domain. LLMs have also been used for multi-agent planning through code generation [16] and task allocation through dialogue and path planning [17]. In this paper, we use LLMs to infer commonsense ways to decompose a goal, but leave formal planning to a PDDL system to ensure plan correctness.

II. BACKGROUND

We formally define the planning problem for single and multi-agent settings. We also provide an overview of PDDL functioning.

A. Classical planning problem

Planning problem P is formally defined as a tuple,

$$\langle \mathcal{S}, s^i, \mathcal{S}^g, \mathcal{A}, \mathcal{O}, \mathcal{P}, \mathcal{T} \rangle$$
,

where: \mathscr{S} is a finite and discrete set of environment states; $s^i \in \mathscr{S}$ is the initial state; $\mathscr{S}^g \subset \mathscr{S}$ is the set of goal states that satisfy a goal conditions g; and $\mathscr{A}, \mathscr{O}, \mathscr{P}$ are sets of symbolic actions, objects, and environment predicates, respectively. A state is defined as a list of environment predicates, applied on objects and agents, that hold true. \mathscr{T} is the underlying transition function that defines how the state conditions change when an action $a \in \mathscr{A}$ is executed in the environment. A solution to P is a symbolic plan π given as a sequence of actions $a_t|_{t=1-T}$, such that preconditions of each a_t holds in the state s_{t-1} , where $s_0 = s^t$ and $s_N \in \mathscr{S}^g$.

```
(:action pickup
:parameters (?ob)
(define (domain blocksworld-4ops)
                                                   precondition (and (clear ?ob)
(on-table ?ob)
(:requirements :strips)
(:predicates (clear ?x)
                                                               (arm-empty))
                                                   :effect (and (holding ?ob)
(not (clear ?ob))
(not (on-table ?ob))
                 (on-table ?x)
                 (arm-empty)
                (holding ?x)
                                                                    (not (arm-empty)))
                 (on ?x ?y))
(a) PDDL domain description
                                                   (b) PDDL action operator
    (define (problem BW-rand-3)
      :objects b1 b2 b3 )
                                                         (unstack b1 b3)
(putdown b1)
    (arm-empty)
                                                          (unstack b3 b2)
     (on-table b1)
                                (and
                                                          stack b3 b1)
    (on b2 b3)
                                (on b3 b2)
                                                         (pickup b2)
    (on b3 b1)
(clear b2)
                                                         (stack b2 b3)
                                                         cost = 6 (unit cost)
            (c) PDDL problem
                                                            (d) PDDL plan
```

Fig. 1: PDDL overview. PDDL planning requires a domain definition, with state predicates (a) and action operators (b), and a problem specification, as an instance of the domain specifying goal conditions (c). A planner can then use these definitions to generate a plan (d).

B. Planning Domain Definition Language (PDDL)

Planning Domain Definition Language (PDDL) serves as a standardized encoding language to represent planning problems. In this paper, we consider deterministic, fullyobservable planning problems. The planning problem in PDDL is represented by two files: domain and problem. The domain file defines the environment in terms of the object types and predicates (Figure 1a), followed by action and their parameters, required preconditions, and postcondition effects (Figure 1b). Some conditions relate to the agent (holding ?x), and others to the environment (on-table ?x). The problem file defines an instance of the domain with a list of objects, initial state and goal conditions (Figure 1c). We use a PDDL planner that takes in a domain and problem file and searches for action sequences leading from the initial state to a goal state, heuristically returning the shortest such plan found, if any (Figure 1d).

C. Multi-agent planning problem

In multi-agent settings, actions are executed by the multiple agents in parallel. In this case, agents share the same environment, so the environment-specific state conditions are shared between the agents. However, agent-specific state conditions are unique to each agent. Agents must consider environment conditions, their own conditions, and other agent's conditions to execute a coordinated action. A multiagent plan consists of actions for each agent at every time step t, $a_t^r|_{t=1-T}^{r=1-|R|}$, for $r \in R$ agents.

III. MULTI-AGENT PLANNING METHOD: TWOSTEP

We propose decomposing multi-agent planning problem into two single agent planning problems by leveraging LLM's commonsense and reasoning contextualized with humans operating together in diverse scenarios. In particular, we consider a two-agent planning problem with a helper and main agent. For a given problem P, the helper plans

 $\pi_h = \prod(i,g')$ to reach a subgoal state g' from P's initial state i, using a planner \prod . This subgoal state (g') then serves as the initial state $(i' = E(i, \prod(s,g')))$ for the main agent plan $\pi_m = \prod(i',g)$ to reach the P's specified goal, g, where E refers to Plan Execution. The helper conceives g' using 2 modules: subgoal generator to produce a possible helper subgoal in English - English subgoal, and subgoal translator to translate English subgoal into PDDL format goal - PDDL subgoal. We hypothesize that the commonsense reasoning abilities of LLMs can be leveraged in this context to predict helper agent subgoals whose plan actions π_h will be executable in parallel to those of the main agent π_m , where latter takes for granted that the former will achieve the subgoal eventually.

To evaluate that hypothesis, we present the TWOSTEP (Figure 2) method.

a) English subgoal generation for helper agent: Given an in-context English subgoal generation example, and the new problem described in English, the LLM generates an English subgoal for this problem. For the in-context English subgoal generation example, we provide the LLM with an example domain containing a textual description of a problem, initial state, and goal conditions, along with the corresponding single agent PDDL plan in this domain. We then append a new problem in this example domain for which the single agent PDDL plan is not provided. Now we provide a potential English subgoal for the helper agent along with the reasoning, such that it reduces some of the goal conditions and eventually the plan length for the main agent. This English subgoal uses hints about the domain functioning from the example single agent plan, but is situated in this new problem, different from the problem for which the plan is provided. The reasoning provided encourages the helper to act independent of the main agent, such that it can complete this subgoal alongside. The reasoning also encourages release of all the objects at the end of its subgoal so that it doesn't block the main agent from using those objects. For instance, in our example in Figure 2, if the helper ends with holding the pot after boiling water, then it will block the main agent from using it to boil the potato slices. After the above English subgoal example, we provide the LLM with the current domain description, different from example domain, specifying an example problem and corresponding single agent PDDL plan, and finally the current problem. This prompting mechanism encourages the LLM to use a similar reasoning process as above to generate an English subgoal in this new problem domain.

b) Translation from English description of subgoals to PDDL goal specifications: To translate the English subgoal into PDDL subgoal, we provide the LLM with the current problem description and initial state in PDDL. We then provide English description of the goal g and its translated PDDL goal as the example, followed by generated English subgoal to generate the corresponding translated PDDL subgoal or g'. We choose the in-context example to be from the current problem as it would encourage the translated PDDL subgoal to be situated in this problem instance, making it to




Fig. 2: Planning pipeline for SA PDDL, MA PDDL, and TWOSTEP. TWOSTEP has a helper and a main agent. The helper tries to extract and complete a partially independent subgoal that reduces the steps for main agent, while main agent completes the remaining task. The plan generated for both agents are then executed to find the execution length.

be parsable and executable by the PDDL planner.

- c) Editing initial state for main agent: We extract changes in the environment state caused by the helper plan π_h to compute the subgoal state g'. We then edit the environment-specific state conditions starting for the original problem's i. These state conditions are not only helper's PDDL subgoal conditions, but also the changes made while executing each step of π_h , to ensure that the interfacing state g' or i' between the two agents are perfectly connected. We do not edit the agent-specific conditions in i so that the main agent can start from its own initial state in a partially solved environment-state and fulfill the conditions remaining for P's original goal g.
- d) Multi-agent plan execution: While the main agent's plan starts at helper agent's end environment state, TWOSTEP encourages the LLM to generate a helper subgoal that can be executed in parallel with the main agent. However, the main agnet may have to have wait for helper to finish taking a certain action during its own plan execution. The Plan Execution maintains common environment state, initialized from the original problem, with individual agent states to run parallel execution for effective execution length computation. We show in our results that the generated subgoal executes in parallel partially, leading to an effectively smaller plan length.

IV. EXPERIMENT SETUP

Our principal experiments measure whether TWOSTEP can lead to more efficient two-agent tasks than a solution to a

single-agent PDDL problem in terms of agent actions while simultaneously requiring less planning time than needed to solve a two-agent PDDL problem formulation. Overall, we find that TWOSTEP achieves that balance when both: the planning domain considers agent state in action preconditions; and the planning problem to solve contains two or more partially independent subgoals rather than requiring strict action sequencing.

A. Domains

We evaluate TWOSTEP on 5 symbolic domains utilized in prior work that leveraged LLMs for planning [15] and in one simulation environment [18].

1) Symbolic Domains: Each symbolic environment defines 20 tasks that vary in number of the environment objects and optimal plan length. These domains are frequently used in classical planning competitions. 1) BARMAN: A bartender with 2 hands has been assigned the job of preparing cocktails for a customer's order, using the specified ingredient and appropriate tools, and striving to accomplish this with the fewest possible actions. 2) BLOCKSWORLD: A collection of block piles arranged on a table, and a robot is tasked to reorganize them into a specified configuration while adhering to the simple physics principles. 3) TERMES: The objective of a robot is to construct complex structures by transporting and positioning blocks, as well as using them as a means to ascend and move adjacent blocks. 4) TYREWORLD: The robot is assigned with changing flat tires, which involves tasks such as removing flat tires, inflating the intact tires, tightening nuts, and returning tools to the boot, all in the correct sequence. This domain does not have a agent state while taking an action, for example, the agent need not be holding a tool first in order to apply a tool to a tire. 5) GRIPPERS: A set of robots with 2 grippers each are given a task to move objects among different rooms.

2) Simulated Domain: We use AI2THOR simulator for multi-agent task execution. We build on the textworld abstraction of AI2THOR, AlfWorld, that enables high-level planning while abstracting low-level execution. It provides the domain and tasks encoded in PDDL for automated task plan generation via classical (e.g. PDDL) planners. We extend the support to multi-agent PDDL planning, by initializing multiple agents in the PDDL initial environment state as well as the simulator. The multi-agent simulation, AI2THOR, is implemented asynchronously, i.e., each agent gets to take an action at every time step, however, due to the underlying AI2THOR support in each simulator timestep only one agent can be issued a command like move forward/pickup/etc. We build a parallel plan execution framework that allows action queuing if an agent is occupied with a previous action, such as 'go to microwave 1' which involves multiple navigation steps. The oracle multi-agent navigation pipeline uses object location and position and other agent's position to navigate. Reachable positions, where an agent can navigate to, at any time change when another agent moves in the environment as it occupies space, akin, to any other object. The simulation supports path replanning around another agent while navigating, however, there are scenarios where the agents are deadlocked together if there's no way around the collision. At this point, we defer a randomly chosen agent to a randomly chosen location to relieve from the deadlock. Thereafter, the agents continue to execute their queued actions.

B. Approaches Considered

We compare TWOSTEP against single agent and multiagent PDDL planning for all the above domains. To enable multi-agent planning for symbolic domains, we edited PDDL domain and problem files to reflect multiple agents. We added 'agent' parameter to all the actions and agent-specific predicates in the domain, and initialized 2 agents along with their initial states in the prolem files of 20 tasks. This extension requires understanding of PDDL functioning. Moreover, initializing another agent increases the number of possible actions at any state as well introduces addition state conditions, leading to larger search space. We also define 4 long-horizon tasks and their corresponding PDDL problem files for our evaluations in the simulated domain, as shown in Table III. Some objects may also be inside another containers like drawers/cabinetes as per the PDDL initial state specification, making the resultant plans longer. We will open-source our defined domain and problem files. We consider the following approaches for the above domains,

1) SA PDDL: PDDL planning in the above domains for a single agent. This baseline will compare TWOSTEP with

a single agent completing the task using a single agent plan generated by the PDDL planner.

- 2) MA PDDL: PDDL planning in the above domains for two agents. This baseline will compare TWOSTEP with multiple agents completing the task using a multi-agent plan generated by the PDDL planner.
- 3) TWOSTEP: multi-agent plan with an LLM-guided helper and main agents our method as shown in Figure 2. This method compares LLM-based division of the task into a partially independent subgoal, executed in parallel, when planning for the two agents are done separately as single agent PDDL problems.

Algorithm 1 Parallel Execution Length Computation Algorithm for symbolic domains

```
1: function EXECUTE(i, j, agent)
        if agent = helper then
 2:
             V_a \leftarrow \text{execute } \pi_h[i]
 3:
             if V_a then
 4:
                 return 1 + \min(
 5:
                       Execute(i+1, j, agent = helper),
 6:
                       Execute(i, j+1, agent = main),
 7:
                       Execute(i+1, j+1, agent = both)
 8:
 9:
             else
                 return 1 \times 10^{10}
10:
             end if
11:
         end if
12:
        if agent = main then
13:
             V_b \leftarrow \text{execute } \pi_m[j]
14:
             if V_b then
15:
                                              ⊳ Same as Line 5-10
16:
             end if
17:
         end if
18:
         if agent = both then
19:
             V_{ab} \leftarrow \text{execute } \pi_h[i], \ \pi_m[j]
20:
21:
             V_{ba} \leftarrow \text{execute } \pi_m[j], \ \pi_h[i]
             if V_{ab} and V_{ba} then
22:
                                              ⊳ Same as Line 5-10
23:
             end if
24:
         end if
26: end function
```

C. Evaluation metrics

We measure planning time and plan execution length to asses the efficacy of plans generated by different approaches described above. A method that achieves lower planning time and a shorter execution length is desirable as it would be faster in finding a plan and executing it. Note that we do not report task success metric, since classical planning methods are always successful. Moreover, TWOSTEP ensures success as the system defaults to a single agent if the inferred subgoal is faulty.

1) Planning Time: total time taken by the PDDL planner to explore the full state space, capped by a time limit. For TWOSTEP, planning time is the sum of helper PDDL

TABLE I: Execution length and planning time in symbolic domains averaged across 20 tasks each. TWOSTEP achieves shorter plans than a single agent without incurring the planning time cost of solving a multi-agent PDDL problem. TWOSTEP execution lengths are highlighted in blue when they average less than SA PDDL and red otherwise. Similarly, TWOSTEP planning times are highlighted in blue when they are less than MA PDDL times.

Execution Length (steps) SA PDDL MA PDDL TWOSTEP	BARMAN	BLOCKSWORLD	TERMES	TYREWORLD	GRIPPERS	Macro Avg
	64.9	17.7	92.2	121.4	8.6	60.9
	41.0	13.4	104.1	121.4	7.0	57.4
	64.0 ± 0.2	16.8 ± 0.2	96.2 ± 0.6	112.3 ± 1.1	6.7±0.1	59.3 ± 0.4
Planning Time (seconds) SA PDDL MA PDDL TWOSTEP	BARMAN	BLOCKSWORLD	TERMES	TYREWORLD	GRIPPERS	Macro Avg
	803.0	273.0	819.0	950.0	1.5	569.3
	1000.0	314.0	966.0	950.0	87.2	663.4
	800.5 ± 6.1	290.0 ± 4.7	797.5 ± 21.3	929.3 ± 172.2	1.4±0.0	563.7 ± 40.9

planning time, main PDDL planning time, and LLM inference time. We report LLM inference time separately as an overall average instead of adding each task's inference time to average out any LLM API call delays.

2) Execution Length: For multi-agent PDDL and TWOSTEP, it refers to the effective length of the plan when multiple agent plans are executed in parallel. For single agent, execution lenth is equal to best PDDL plan cost, however, for multi-agent, execution lenth ≤ best PDDL plan cost. For symbolic domains, we utilize Algorithm 1 to find the best possible execution length, by dynamically finding the best parallel plan execution strategy. For simulated domain, execution length refers to number of environment steps taken, where all agents take a step sequentially in each environment step, including navigation and interaction steps. Depending on action preconditions and execution state, we also consider some agent's action at a time step to be a 'wait' or a 'no operation' action.

D. Experiment Details

PDDL planning for all approaches use FAST-DOWNWARD planner¹ with LAMA alias and a search time limit of 1000 seconds. For plan validation and extracting state condition changes with the plan execution, we use VAL². For TwoStep's LLM-based modules in Figure 2, we use GPT-4 provided as an API from OpenAI³. We use the top probability response from the LLM. The results are reported for an average of 3 runs with standard deviations for TwoStep as it involves LLM inference. We found that the FAST-DOWNWARD doesn't add any variance to the pipeline, therefore we only report 1 run for all other methods as they don't involve LLM predictions.

V. RESULTS

We report results for the above defined systems — SA PDDL, MA PDDL, TWOSTEP — in symbolic and simulated domains. We also report an ablation of TWOSTEP by replacing helper's LLM-based subgoal inference by a human in the symbolic domains. We report averaged execution lengths and planning times in Table I on each of the 5 domains with 20 problems each.

TWOSTEP takes comparable planning time to that of SA PDDL, to come up with a more efficient multiagent plan for 4/5 domains, while taking notably less planning time than MA PDDL. The longer time taken in TYREWORLD is because sometimes the helper agent ends up searching the state space for longer due to generating a bigger subgoal to achieve. LLM inference times are 6.79 \pm 1.48 for English subgoal generation and 4.67 \pm 2.55 for PDDL subgoal generation. The LLM inference times depend on the size of input and generated tokens, which may correspond to the size of the domain. The monetary cost of GPT-4 inference was \$30.

The execution length for TWOSTEP is slightly shorter for 3/5 domains, and significantly shorter for TYREWORLD, compared to that of SA PDDL. We observe that MA PDDL is able to find a significantly shorter plan for those 3/5 domains. This is perhaps because MA PDDL is able to exhaust the state space search and is able to come up with shorter plan, which gets further shorter when parallelizing. Moreover, these domains come with defined agent-specific state conditions such as agent's inventory or location, which leads to a plan where multiple agents can work with different objects so they don't have to empty their hands before interacting with another object, or can work at different locations as well. However, MA PDDL is not trying to optimize for parallelizability of the plan, it happens as a byproduct of these agent-specific state conditions. When these state conditions are missing, such as in TYREWORLD, MA PDDL is only as good as single agent PDDL. Moreover, when the domain is very complex, even with these conditions specified, MA PDDL can't optimize the plan length within the specified time limit, as we see for TERMES, where it gets worse than SA PDDL, as adding another agent significantly increase the domain complexity. While TWOSTEP has lower execution length than MA PDDL, but is still worse than that of SA PDDL in TERMES domain.

TWOSTEP works with the SA PDDL framework as it is without needing to change the PDDL definitions, and works with or without agent-specific conditions specified by leveraging LLM commonsense. However, MA PDDL planning requires domain and problem definitions to be multi-agent compatible, and heavily depends on whether agent-specific conditions are defined.

¹https://github.com/aibasel/downward/tree/release-22.12.0

²https://github.com/KCL-Planning/VAL

³https://platform.openai.com/docs/models/gpt-4

TABLE II: Execution length in 4 symbolic domains averaged across 4 tasks each. TWOSTEP achieves comparable execution lengths as an expert human annotated subgoals (TWOSTEP-Human).

Execution Length (steps)	BLOCKSWORLD	TERMES	TYREWORLD	GRIPPERS	Macro Avg
TWOSTEP	15.5 ± 0.4	125.5 ± 0.8	88.3 ± 1.9	6.5 ± 0.0	58.9 ± 0.8
TWOSTEP-Human	16.3	120.0	86.8	5.3	57.1

TABLE III: Execution length and planning time in simulated environment averaged across 4 tasks each. On average, TWOSTEP achieves comparable length plans as a single agent without incurring the planning time cost of solving a large goal or a multi-agent PDDL problem. TWOSTEP execution lengths are highlighted in blue when they average less than SA PDDL and red otherwise. Similarly, TWOSTEP planning times are highlighted in blue when they are less than MA PDDL times. timeout indicates reaching a time limit of 60 minutes.

Tasks	 Put an apple, a heated egg, and a clean fork on the dining table Clear out the dining table by removing the plate, two bowls, and the mug Put the apple, the egg, and the wine bottle in the fridge Clear out the sink of the spoon, the cup, the knife, and the apple 					
Execution Length (steps) SA PDDL MA PDDL TWOSTEP	Task 1 172.0 189.2 ± 16.7 113.6 ± 6.6	Task 2 201.0 260.0 ± 248.5 147.0 ± 62.7	Task 3 99.0 196.7 ± 88.3 204.5 ± 17.8	Task 4 44.0 166.5 ± 69.6 54.2 ± 14.0	Macro Avg 129.0 203.1 ± 105.8 129.8 ± 25.3	
Planning Time (seconds) SA PDDL MA PDDL TWOSTEP	Task 1 timeout timeout 22.0	Task 2 timeout timeout 14.8	Task 3 11.3 timeout 18.4	Task 4 timeout timeout timeout	3 timeouts 4 timeouts 1 timeout	

LLM subgoal generation approximates human expert subgoal specification. We perform human studies to assess the quality LLM inferred subgoals in TWOSTEP. In the first study, we ask a PDDL expert (a graduate student working with PDDL) to annotate the subgoal in PDDL goal format for 4 domains — BLOCKSWORLD, TERMES, TYREWORLD, GRIPPERS — 4 tasks each, sampled uniformly between tasks 1-20. It is referred to as TWOSTEP-Human as we replace LLM subgoal with human subgoal in TWOSTEP. The expert human took 70 minutes to annotate 16 tasks. We report execution lengths in Table II. It shows that TWOSTEP performance quite close to an expert written subgoals. We also ask human unfamiliar with PDDL to annotate these tasks with natural language subgoals, however, a non-expert human found it quite challenging to write subgoals for these symbolic domains.

TwoStep advantages generalization to temporal simulation setting without temporal awareness. We report AI2THOR planning time results in Table III and and execution length in Table III for the 3 approaches. Planning time results show that both SA PDDL and MA PDDL take a long time to plan and don't converge quickly. SA PDDL has a large goal with several goal conditions to achieve, and MA PDDL additionally has added state space complexity due to an additional agent. In execution length, SA PDDL has no variance. MA PDDL and TWOSTEP have variance in length due to random initialization of agent start location as well as random deadlock release methodology. The deadlocks also cause TWOSTEP to be not achieve lower execution length than SA PDDL as they add additional coordination steps. Moreover, MA PDDL suffers a lot more with this problem, as it tries to utilize both agents back and forth, while TWOSTEP only splits the problem with one subgoal.

VI. CONCLUSION

We propose TWOSTEP, a method to decompose a single agent planning problem into a two-agent planning problem in several symbolic domains and one simulated domain. TWOSTEP leverages commonsense from LLM to effectively divide a problem between two agents for faster execution, while also preserving execution success using classical planning guarantees. Our results show that LLM-based goal decomposition leads to faster planning time than the multiagent PDDL problem and shorter plan execution steps than the single agent execution. We additionally show that LLM-inferred subgoals in TWOSTEP approximate those specified by a human expert.

ACKNOWLEDGEMENTS

This work was supported by a grant from the Army Research Lab (ARL) Army AI Innovations Institute (A2I2), award number W911NF-23-2-0010.

REFERENCES

- [1] R. E. Fikes and N. J. Nilsson, "Strips: A new approach to the application of theorem proving to problem solving," *Artificial Intelligence*, vol. 2, no. 3, pp. 189–208, 1971. [Online]. Available: https://www.sciencedirect.com/science/article/pii/0004370271900105
- [2] J. Carbonell, O. Etzioni, Y. Gil, R. Joseph, C. Knoblock, S. Minton, and M. Veloso, "Prodigy: An integrated architecture for planning and learning," SIGART Bull., vol. 2, no. 4, p. 51–55, jul 1991. [Online]. Available: https://doi.org/10.1145/122344.122353
- [3] D. Nau, T.-C. Au, O. Ilghami, U. Kuter, J. W. Murdock, D. Wu, and F. Yaman, "Shop2: An htn planning system," J. Artif. Intell. Res. (JAIR), vol. 20, pp. 379–404, 12 2003.
- [4] Y.-q. Jiang, S.-q. Zhang, P. Khandelwal, and P. Stone, "Task planning in robotics: an empirical comparison of pddl- and asp-based systems," Frontiers of Information Technology & Electronic Engineering, vol. 20, pp. 363–373, 03 2019.
- [5] M. Fox and D. Long, "Pddl2.1: An extension to pddl for expressing temporal planning domains," ArXiv, vol. abs/1106.4561, 2003. [Online]. Available: https://api.semanticscholar.org/CorpusID:1397894

- [6] G. Brewka, T. Eiter, and M. Truszczyński, "Answer set programming at a glance," *Commun. ACM*, vol. 54, no. 12, p. 92–103, dec 2011. [Online]. Available: https://doi.org/10.1145/2043174.2043195
- [7] V. Lifschitz, "Answer set programming and plan generation," Artificial Intelligence, vol. 138, no. 1, pp. 39–54, 2002, knowledge Representation and Logic Programming. [Online]. Available: https://www.sciencedirect.com/science/article/pii/S0004370202001868
- [8] W. Huang, P. Abbeel, D. Pathak, and I. Mordatch, "Language models as zero-shot planners: Extracting actionable knowledge for embodied agents," in *Proceedings of the 39th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, K. Chaudhuri, S. Jegelka, L. Song, C. Szepesvari, G. Niu, and S. Sabato, Eds., vol. 162. PMLR, 17–23 Jul 2022, pp. 9118–9147. [Online]. Available: https://proceedings.mlr.press/v162/huang22a.html
- [9] B. Ichter, A. Brohan, Y. Chebotar, C. Finn, K. Hausman, A. Herzog, D. Ho, J. Ibarz, A. Irpan, E. Jang, R. Julian, D. Kalashnikov, S. Levine, Y. Lu, C. Parada, K. Rao, P. Sermanet, A. T. Toshev, V. Vanhoucke, F. Xia, T. Xiao, P. Xu, M. Yan, N. Brown, M. Ahn, O. Cortes, N. Sievers, C. Tan, S. Xu, D. Reyes, J. Rettinghouse, J. Quiambao, P. Pastor, L. Luu, K.-H. Lee, Y. Kuang, S. Jesmonth, K. Jeffrey, R. J. Ruano, J. Hsu, K. Gopalakrishnan, B. David, A. Zeng, and C. K. Fu, "Do as i can, not as i say: Grounding language in robotic affordances," in 6th Annual Conference on Robot Learning, 2022. [Online]. Available: https://openreview.net/forum?id=bdHkMjBJG_w
- [10] I. Singh, V. Blukis, A. Mousavian, A. Goyal, D. Xu, J. Tremblay, D. Fox, J. Thomason, and A. Garg, "Progprompt: Generating situated robot task plans using large language models," in 2023 IEEE International Conference on Robotics and Automation (ICRA), 2023, pp. 11523–11530.
- [11] J. Liang, W. Huang, F. Xia, P. Xu, K. Hausman, B. Ichter, P. Florence, and A. Zeng, "Code as policies: Language model programs for embodied control," in 2023 IEEE International Conference on Robotics and Automation (ICRA), 2023, pp. 9493–9500.
- [12] I. Singh, V. Blukis, A. Mousavian, A. Goyal, D. Xu, J. Tremblay, D. Fox, J. Thomason, and A. Garg, "Progprompt: program generation for situated robot task planning using large language models," *Autonomous Robots*, 08 2023.
- [13] D. Trivedi, J. Zhang, S.-H. Sun, and J. J. Lim, "Learning to synthesize programs as interpretable and generalizable policies," in Advances in Neural Information Processing Systems, vol. 34, 2021. [Online]. Available: https://arxiv.org/pdf/2108.13643.pdf
- [14] T. Silver, S. Dan, K. Srinivas, J. B. Tenenbaum, L. P. Kaelbling, and M. Katz, "Generalized planning in pddl domains with pretrained large language models," 2023.
- [15] B. Liu, Y. Jiang, X. Zhang, Q. Liu, S. Zhang, J. Biswas, and P. Stone, "Llm+p: Empowering large language models with optimal planning proficiency," arXiv preprint arXiv:2304.11477, 2023.
- [16] S. S. Kannan, V. L. N. Venkatesh, and B.-C. Min, "Smartllm: Smart multi-agent robot task planning using large language models," *ArXiv*, vol. abs/2309.10062, 2023. [Online]. Available: https://api.semanticscholar.org/CorpusID:262055166
- [17] Z. Mandi, S. Jain, and S. Song, "Roco: Dialectic multi-robot collaboration with large language models," 2023.
- [18] E. Kolve, R. Mottaghi, W. Han, E. VanderBilt, L. Weihs, A. Herrasti, D. Gordon, Y. Zhu, A. Gupta, and A. Farhadi, "AI2-THOR: An Interactive 3D Environment for Visual AI," arXiv, 2017.

APPENDIX

A. Subgoal Generator Prompt

System Prompt

Main agent: agent0 Helper agent: agent1

Generate one clearly stated small independent subgoal for helper agent to help main agent complete the given task. The subgoal should not rely on any main agent actions, and should be executable by the helper agent independently without waiting for any main agent actions. The subgoal should be clearly state with unambiguous terminology. Do not use actions like assist or help. Generate actions that the helper agent can do independently, based on the given steps for completing the task. The main agent should be able to continue working on the remaining task while the helper agent is completing its small subgoal. Do not overtake the full sequence of actions. Remember, the helper agent is only assisting the main agent and acts agnostic to the main agent.

User Prompt and LLM Response

Part 1: Fixed

Example domain scenario:

You have 1 shaker with 3 levels, 3 shot glasses, 3 dispensers for 3 ingredients. The shaker and shot glasses are clean, empty, and on the table. Your left and right hands are empty. The first ingredient of cocktail1 is ingredient3. The second ingredient of cocktail1 is ingredient1. Your goal is to make 1 cocktail. shot1 contains cocktail1.

agent0 takes the following steps to complete the above task:

grasp right shot2

10 11

fill-shot shot2 ingredient1 right left dispenser1 pour-shot-to-clean-shaker shot2 ingredient1 shaker1 right

clean-shot shot2 ingredient1 right left fill-shot shot2 ingredient3 right left dispenser3 grasp left shaker1

pour-shot-to-used-shaker shot2 ingredient3 shaker1 right 11 12

leave right shot2

shake cocktail1 ingredient3 ingredient1 shaker1 left right pour-shaker-to-shot cocktail1 shot1 left shaker1 l2 l1

Now we have a new problem defined in this domain for which we don't have access to the signle agent plan: You have 1 shaker with 3 levels, 3 shot glasses, 3 dispensers for 3 ingredients. The shaker and shot glasses are clean, empty, and on the table. Your left and right hands are empty. The first ingredient of cocktail1 is ingredient3. The second ingredient of cocktail2 is ingredient1. The first ingredient of cocktail2 is ingredient1. The second ingredient of cocktail2 is ingredient2. Your goal is to make 2 cocktails. shot1 contains cocktail1. shot2 contains cocktail2.

A possible agent1 subgoal looking at how the domain works based on the plan example provided for another task in this domain could be -

agent1 subgoals: It can help in filling ingredient1 in a shot glass, then pour it in shaker1, while agent0 prepares other cocktail ingredients using other objects. In this way, agent1 would not need to wait for agent0 and it can complete its goal independently. agent1 should also release all objects that the main agent might need for its own actions. Therefore, agent1's clearly stated (with object names) complete and final goal condition is: shaker1 contains ingredient1 and all hands are empty.

Part 2: Domain specific (BLOCKSWORLD)

Current domain scenario:

You have 5 blocks. b2 is on top of b5. b5 is on top of b1. b1 is on top of b4. b3 is on top of b2. b4 is on the table. b3 is clear. Your arm is empty. Your goal is to move the blocks. b4 should be on top of b3.

agent0 takes the following steps to complete the above task:

unstack b3 b2

putdown b3

unstack b2 b5

putdown b2

unstack b5 b1

putdown b5

unstack b1 b4

putdown b1

pickup b4

stack b4 b3

Now we have a new problem defined in this domain for which we don't have access to the signle agent plan: You have 3 blocks. b2 is on top of b3. b3 is on top of b1. b1 is on the table. b2 is clear. Your arm is empty. Your goal is to move the blocks. b3 should be on top of b2. b1 should be on top of b3.

A possible agent1 subgoal looking at how the domain works based on the plan example provided for another task in this domain could be - agent1 subgoals:

LLM response - English Subgoal (BLOCKSWORLD)

It can help in unstacking block b2 from b3 and putting it down on the table, while agent0 works on rearranging the other blocks. In this way, agent1 would not need to wait for agent0 and it can complete its goal independently. agent1 should also release all objects that the main agent might need for its own actions. Therefore, agent1's clearly stated (with object names) complete and final goal condition is: b2 is on the table and the arm is empty.

Part 2: Domain specific (TERMES)

Current domain scenario:

The robot is on a grid with 3 rows and 3 columns.

pos-0-0 pos-0-1 pos-0-2

pos-1-0 pos-1-1 pos-1-2

pos-2-0 pos-2-1 pos-2-2

The robot is at pos-2-0.

The depot for new blocks is at pos-2-0. The maximum height of blocks is 2. Your goal is to build blocks so that the height at pos-1-1 is 2. You cannot have an unplaced block at the end.

agent0 takes the following steps to complete the above task:

create-block pos-2-0

place-block pos-2-0 pos-1-0 n0 n1

create-block pos-2-0

move pos-2-0 pos-2-1 n0

place-block pos-2-1 pos-1-1 n0 n1

move pos-2-1 pos-2-0 n0

create-block pos-2-0

move-up pos-2-0 n0 pos-1-0 n1

place-block pos-1-0 pos-1-1 n1 n2

move-down pos-1-0 n1 pos-2-0 n0

remove-block pos-2-0 pos-1-0 n1 n0

destroy-block pos-2-0

Now we have a new problem defined in this domain for which we don't have access to the signle agent plan:

The robot is on a grid with 4 rows and 3 columns.

pos-0-0 pos-0-1 pos-0-2

pos-1-0 pos-1-1 pos-1-2

pos-2-0 pos-2-1 pos-2-2

pos-3-0 pos-3-1 pos-3-2

The robot is at pos-2-0.

The depot for new blocks is at pos-2-0. The maximum height of blocks is 3. Your goal is to build blocks so that the height at pos-1-2 is 3. You cannot have an unplaced block at the end.

A possible agent1 subgoal looking at how the domain works based on the plan example provided for another task in this domain could be - agent1 subgoals:

LLM response - English Subgoal (TERMES)

It can help in creating a block at the depot pos-2-0, then move it to pos-1-2, while agent0 prepares other blocks using other positions. In this way, agent1 would not need to wait for agent0 and it can complete its goal independently. agent1 should also ensure that there are no unplaced blocks at the end. Therefore, agent1's clearly stated (with object names) complete and final goal condition is: pos-1-2 has a block of height n1 and all unplaced blocks are destroyed.

Part 2: Domain specific (TYREWORLD)

Current domain scenario:

You have a jack, a pump, a wrench, a boot, 2 hubs, 2 nuts, 2 flat tyres, and 2 intact tyres. The jack, pump, wrench, and intact tyres are in the boot. The boot is unlocked but is closed. The intact tyres are not inflated. The flat tyres are on the hubs. The hubs are on the ground. The nuts are tight on the hubs. The hubs are fastened. Your goal is to replace flat tyres with intact tyres on the hubs. Intact tyres should be inflated. The nuts should be tight on the hubs. The flat tyres, wrench, jack, and pump should be in the boot. The boot should be closed.

agent0 takes the following steps to complete the above task: loosen nuts2 the-hub2 jack-up the-hub2 undo nuts2 the-hub2 remove-wheel w2 the-hub2 loosen nuts1 the-hub1 jack-up the-hub1 undo nuts1 the-hub1 remove-wheel w1 the-hub1 inflate r2 inflate r1 open boot fetch r1 boot put-on-wheel r1 the-hub1 do-up nuts1 the-hub1 jack-down the-hub1 tighten nuts1 the-hub1 fetch r2 boot put-on-wheel r2 the-hub2 do-up nuts2 the-hub2 jack-down the-hub2 tighten nuts2 the-hub2 put-away w1 boot put-away w2 boot close boot

Now we have a new problem defined in this domain for which we don't have access to the signle agent plan: You have a jack, a pump, a wrench, a boot, 1 hubs, 1 nuts, 1 flat tyres, and 1 intact tyres. The jack, pump, wrench, and intact tyres are in the boot. The boot is unlocked but is closed. The intact tyres are not inflated. The flat tyres are on the hubs. The hubs are on the ground. The nuts are tight on the hubs. The hubs are fastened. Your goal is to replace flat tyres with intact tyres on the hubs. Intact tyres should be inflated. The nuts should be tight on the hubs. The flat tyres, wrench, jack, and pump should be in the boot. The boot should be closed.

A possible agent1 subgoal looking at how the domain works based on the plan example provided for another task in this domain could be - agent1 subgoals:

LLM response - English Subgoal (TYREWORLD)

It can help in inflating the intact tyre while agent0 works on removing the flat tyre from the hub. In this way, agent1 would not need to wait for agent0 and it can complete its goal independently. agent1 should also release all objects that the main agent might need for its own actions. Therefore, agent1's clearly stated (with object names) complete and final goal condition is: the intact tyre is inflated and all hands are empty.

Part 2: Domain specific (GRIPPERS)

Current domain scenario:

You control 1 robot, each robot has a left gripper and a right gripper. There are 4 rooms and 6 balls. robot1 is in room3. ball1 is in room3. ball2 is in room1. ball3 is in room3. ball4 is in room2. ball5 is in room4. ball6 is in room4. The robots' grippers are free. Your goal is to transport the balls to their destinations. ball1 should be in room1. ball4 should be in room2. ball5 should be in room1. ball6 should be in room1.

agent0 takes the following steps to complete the above task:
pick robot1 ball1 room3 lgripper2
move robot1 room3 room1
move robot1 room1 room4
pick robot1 ball5 room4 rgripper2
drop robot1 ball6 room4 lgripper2
pick robot1 ball6 room4 lgripper2
move robot1 room4 room1
drop robot1 ball6 room1 lgripper2
drop robot1 ball5 room1 rgripper2
move robot1 room3 room3
pick robot1 ball3 room3 lgripper2
move robot1 room3 room1
drop robot1 ball3 room1 lgripper2

Now we have a new problem defined in this domain for which we don't have access to the signle agent plan: You control 1 robots, each robot has a left gripper and a right gripper. There are 2 rooms and 2 balls. robot1 is in room1. ball2 is in room1. ball1 is in room1. The robots' grippers are free. Your goal is to transport the balls to their destinations. ball1 should be in room2. ball2 should be in room2.

A possible agent1 subgoal looking at how the domain works based on the plan example provided for another task in this domain could be - agent1 subgoals:

LLM response - English Subgoal (GRIPPERS)

It can help in picking up ball1 in room1 and transporting it to room2. In this way, agent1 would not need to wait for agent0 and it can complete its goal independently. agent1 should also release all objects that the main agent might need for its own actions. Therefore, agent1's clearly stated (with object names) complete and final goal condition is: ball1 is in room2 and all grippers are free.

B. Subgoal Translator Prompt

System Prompt

I want you to solve planning problems. Provide me with the PDDL goal that describes the new planning goal directly without further explanations. Make sure to provide only non-conflicting, necessary, and final goal conditions mentioned in the given goal.

User Prompt and LLM Response

```
Domain and Problem specific (BLOCKSWORLD)
```

```
The PDDL problem and its initial conditions are given as:
(define (problem BW-rand-3)
(:domain blocksworld-4ops)
(:objects b1 b2 b3 )
(:init
(arm-empty)
(on-table b1)
(on b2 b3)
(on b3 b1)
(clear b2)
)
```

An example planning goal for this problem: Your goal is to move the blocks. b3 should be on top of b2. b1 should be on top of b3.

```
The PDDL goal for the example planning goal: (:goal (and (on b3 b2) (on b1 b3))
```

New planning goal for the same problem: Your goal is: b2 is on the table and the arm is empty.

The PDDL goal for the new planning goal:

```
LLM response - PDDL Subgoal (BLOCKSWORLD)
(:goal
(and
(on-table b2)
(arm-empty))
)
```

```
Domain and Problem specific (TERMES)
 The PDDL problem and its initial conditions are given
 as:
                             termes-00038-0036-4x3x3-
 (define
              (problem
 random_towers_4x3_3_1_3)
 (:domain termes)
 ; termes-00038-0036-4x3x3-random_towers_4x3_3_1_3
 ; Initial state:
 ; 0 0 R0D 0
 ; 0 0 0 0
 ;0000
 ; Goal state:
 ; 0 0 0 0
 ;0000
 ; 0 3 0 0
 ; Maximal height: 3
 (:objects
 n0 - numb
 n1 - numb
 n2 - numb
 n3 - numb
 pos-0-0 - position
 pos-0-1 - position
 pos-3-0 - position
 pos-3-1 - position
 pos-3-2 - position
 (:init
 (height pos-0-0 n0)
 (height pos-3-0 n0)
 (height pos-3-1 n0)
 (height pos-3-2 n0)
 (at pos-2-0)
 (SUCC n1 n0)
 (SUCC n2 n1)
 (SUCC n3 n2)
 (NEIGHBOR pos-0-0 pos-1-0)
 (NEIGHBOR pos-0-0 pos-0-1)
 (NEIGHBOR pos-0-1 pos-1-1)
 (NEIGHBOR pos-0-1 pos-0-0)
 (NEIGHBOR pos-0-1 pos-0-2)
 (NEIGHBOR pos-0-2 pos-1-2)
 (NEIGHBOR pos-3-2 pos-2-2)
 (NEIGHBOR pos-3-2 pos-3-1)
 (IS-DEPOT pos-2-0)
 An example planning goal for this problem:
 Your goal is to build blocks so that the height at pos-1-2
 is 3. You cannot have an unplaced block at the end.
```

```
The PDDL goal for the example planning goal:
 (:goal
 (and
 (height pos-0-0 n0)
 (height pos-0-1 n0)
 (height pos-0-2 n0)
 (height pos-1-0 n0)
 (height pos-1-1 n0)
 (height pos-1-2 n3)
 (height pos-2-0 n0)
 (height pos-2-1 n0)
 (height pos-2-2 n0)
 (height pos-3-0 n0)
 (height pos-3-1 n0)
 (height pos-3-2 n0)
 (not (has-block))
 ))
 New planning goal for the same problem:
 Your goal is: pos-1-2 has a block of height n1 and the
 robot is at pos-2-0.
 The PDDL goal for the new planning goal:
LLM response - PDDL Subgoal (TERMES)
 (:goal
 (and
 (height pos-0-0 n0)
 (height pos-0-1 n0)
```

```
(height pos-0-2 n0)
(height pos-1-0 n0)
(height pos-1-1 n0)
(height pos-1-2 n1)
(height pos-2-0 n0)
(height pos-2-1 n0)
(height pos-2-2 n0)
(height pos-3-0 n0)
(height pos-3-1 n0)
(height pos-3-2 n0)
(not (has-block))
))
```

```
Domain and Problem specific (TYREWORLD)
 The PDDL problem and its initial conditions are given
 (define (problem tyreworld-1)
 (:domain tyreworld)
 (:objects
 wrench jack pump - tool
 the-hub1 - hub
 nuts1 - nut
 boot - container
 r1 w1 - wheel
 (:init
 (in jack boot)
 (in pump boot)
 (in wrench boot)
 (unlocked boot)
 (closed boot)
 (intact r1)
 (in r1 boot)
 (not-inflated r1)
 (on w1 the-hub1)
 (on-ground the-hub1)
 (tight nuts1 the-hub1)
 (fastened the-hub1)
 An example planning goal for this problem:
 Your goal is to replace flat tyres with intact tyres on the
 hubs. Intact tyres should be inflated. The nuts should be
 tight on the hubs. The flat tyres, wrench, jack, and pump
 should be in the boot. The boot should be closed.
 The PDDL goal for the example planning goal:
 (:goal
 (and
 (on r1 the-hub1)
 (inflated r1)
```

```
The PDDL goal for the example planning goal:
(:goal
(and
(on r1 the-hub1)
(inflated r1)
(tight nuts1 the-hub1)
(in w1 boot)
(in wrench boot)
(in jack boot)
(in pump boot)
(closed boot)
))

New planning goal for the same problem:
Your goal is: the intact tyre is inflated and the pump is in the boot.
```

The PDDL goal for the new planning goal:

```
(:goal
 (and
 (inflated r1)
 )
Domain and Problem specific (GRIPPERS)
 The PDDL problem and its initial conditions are given
 (define (problem gripper-2-2-2)
 (:domain gripper-strips)
 (:objects robot1 - robot
 rgripper1 lgripper1 - gripper
 room1 room2 - room
 ball1 ball2 - object)
 (:init
 (at-robby robot1 room1)
 (free robot1 rgripper1)
 (free robot1 lgripper1)
 (at ball1 room1)
 (at ball2 room1)
 An example planning goal for this problem:
 Your goal is to transport the balls to their destinations.
 ball1 should be in room2. ball2 should be in room2.
 The PDDL goal for the example planning goal:
 (:goal
 (and
 (at ball1 room2)
 (at ball2 room2)
 ))
 New planning goal for the same problem:
 Your goal is: ball1 is in room2 and all grippers are free.
 The PDDL goal for the new planning goal:
LLM response - PDDL Subgoal (GRIPPERS)
 (:goal
 (and
 (at ball1 room2)
 (free robot1 rgripper1)
 (free robot1 lgripper1)
```

LLM response - PDDL Subgoal (TYREWORLD)